

```
In [0]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import sqlite3
from sqlalchemy import create_engine # database connection
import csv
import os
warnings.filterwarnings("ignore")
import datetime as dt
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
#from sklearn.cross_validation import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve
from sklearn.model_selection import train_test_split
from scipy.sparse import hstack
from sklearn.preprocessing import StandardScaler
```

```
In [0]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
In [0]: #prepro_features_train.csv (Simple Preprocessing Features)
#nlp_features_train.csv (NLP Features)
if os.path.isfile("drive/My Drive/Quora/nlp_features_train.csv"):
    dfnlp = pd.read_csv("drive/My Drive/Quora/nlp_features_train.csv", encoding='latin-1')
else:
    print("download nlp_features_train.csv from drive or run previous notebook")

if os.path.isfile("drive/My Drive/Quora/df_fe_without_preprocessing_train.csv"):
    dfppro = pd.read_csv("drive/My Drive/Quora/df_fe_without_preprocessing_train.csv", encoding='latin-1')
else:
    print("download df_fe_without_preprocessing_train.csv from drive or run previous notebook")
```

```
In [0]: dfnlp.columns
```

```
Out[0]: Index(['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate',
             'cwc_min', 'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',
             'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',
             'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
             'fuzz_partial_ratio', 'longest_substr_ratio'],
            dtype='object')
```

```
In [0]: df1 = dfnlp.drop(['qid1', 'qid2', 'question1', 'question2'], axis=1) #it has all advanced features
df2 = dfppro.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1) #it has all basic features
#df3 = df.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1)
#df3_q1 = pd.DataFrame(df3.q1_feats_m.values.tolist(), index= df3.index)
#df3_q2 = pd.DataFrame(df3.q2_feats_m.values.tolist(), index= df3.index)
df3=dfnlp[['id', 'question1', 'question2']]
df3['question1'] = df3['question1'].apply(lambda x: str(x))
df3['question2'] = df3['question2'].apply(lambda x: str(x))
```

```
In [0]: print(df1.columns)
print(df2.columns)
print(df3.columns)

Index(['id', 'is_duplicate', 'cwc_min', 'cwc_max', 'csc_min', 'csc_max',
      'ctc_min', 'ctc_max', 'last_word_eq', 'first_word_eq', 'abs_len_diff',
      'mean_len', 'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',
      'fuzz_partial_ratio', 'longest_substr_ratio'],
      dtype='object')
Index(['id', 'freq_qid1', 'freq_qid2', 'qlen', 'q2len', 'q1_n_words',
      'q2_n_words', 'word_Common', 'word_Total', 'word_share', 'freq_q1+q2',
      'freq_q1-q2'],
      dtype='object')
Index(['id', 'question1', 'question2'], dtype='object')
```

```
In [0]: df4=pd.DataFrame()  
df4['id']=df3['id']  
df4['text']=df3['question1']+' '+df3['question2']
```

```
In [0]: df1= df1.merge(df2, on='id',how='left')  
X= df1.merge(df4, on='id',how='left')  
Y=X.is_duplicate  
X=X.drop(['is_duplicate','id'],axis=1)
```

```
In [0]: X.columns
```

```
Out[0]: Index(['cwc_min', 'cwc_max', 'csc_min', 'csc_max', 'ctc_min', 'ctc_max',  
              'last_word_eq', 'first_word_eq', 'abs_len_diff', 'mean_len',  
              'token_set_ratio', 'token_sort_ratio', 'fuzz_ratio',  
              'fuzz_partial_ratio', 'longest_substr_ratio', 'freq_qid1', 'freq_qid2',  
              'q1len', 'q2len', 'q1_n_words', 'q2_n_words', 'word_Common',  
              'word_Total', 'word_share', 'freq_q1+q2', 'freq_q1-q2', 'text'],  
              dtype='object')
```

```
In [0]: #taking only top 100k datapoints  
X=X[0:100000]  
Y=Y[0:100000]
```

```
In [0]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3  
3)  
tf_idf_vect = TfidfVectorizer()  
tf_idf_vect.fit(X_train.text)  
X_train_tf_idf = tf_idf_vect.transform(X_train.text)  
X_test_tf_idf = tf_idf_vect.transform(X_test.text)
```

```
In [0]: X_train=X_train.drop(['text'],axis=1)  
X_test=X_test.drop(['text'],axis=1)
```

```
In [0]: a=np.array(X_train_tf_idf)
```

```
In [0]: print(type(X_train.values))  
print(type(X_train_tf_idf))
```

```
<class 'numpy.ndarray'>  
<class 'scipy.sparse.csr.csr_matrix'>
```

```
In [0]: X_train= hstack((X_train.values,X_train_tf_idf))  
X_test= hstack((X_test.values,X_test_tf_idf))
```

```
In [0]: #below is the function to vectorize questions and returns the stacked x  
train ,xtest and xcv  
def vectorize_and_add_to_df(X,Y):  
    X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=  
0.33) # this is random splitting  
  
    #applying tfidf vectorizer  
    tf_idf_vect = TfidfVectorizer()  
    tf_idf_vect.fit(X_train.text)  
    X_train_tf_idf = tf_idf_vect.transform(X_train.text)  
    X_test_tf_idf = tf_idf_vect.transform(X_test.text)  
  
    #removing Test columns from original dataframe  
    X_train=X_train.drop(['text'],axis=1)  
    X_test=X_test.drop(['text'],axis=1)  
  
    #stacking tfidf features with old features  
    X_train= hstack((X_train.values,X_train_tf_idf))  
    X_test= hstack((X_test.values,X_test_tf_idf))  
  
    scale = StandardScaler(with_mean=False)  
    X_train = scale.fit_transform(X_train)  
    X_test = scale.transform(X_test)  
    return X_train,X_test,y_train,y_test
```

```
In [0]: X_train,X_test,y_train,y_test=vectorize_and_add_to_df(X,Y)
```

```
In [0]: print(X_train.shape)
        print(X_test.shape)
```

```
(67000, 38138)
(33000, 38138)
```

Task 1

confusion matrix

```
In [0]: # This function plots the confusion matrices given y_i, y_i_hat.
        def plot_confusion_matrix(test_y, predict_y):
            C = confusion_matrix(test_y, predict_y)
            # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted class j

            A = (((C.T)/(C.sum(axis=1))).T)
            #divid each element of the confusion matrix with the sum of elements in that column

            # C = [[1, 2],
            #       [3, 4]]
            # C.T = [[1, 3],
            #         [2, 4]]
            # C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to rows in two dimensional array
            # C.sum(axis=1) = [[3, 7]]
            # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
            #                             [2/3, 4/7]]

            # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
            #                               [3/7, 4/7]]
            # sum of row elements = 1

            B = (C/C.sum(axis=0))
            #divid each element of the confusion matrix with the sum of element
```

```

s in that row
    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0) axis=0 corresponds to columns and axis=1 corresponds
    # to rows in two dimensional array
    # C.sum(axis=0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]
    plt.figure(figsize=(20,4))

    labels = [1,2]
    # representing A in heatmap format
    cmap=sns.light_palette("blue")
    plt.subplot(1, 3, 1)
    sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels
, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Confusion matrix")

    plt.subplot(1, 3, 2)
    sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels
, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Precision matrix")

    plt.subplot(1, 3, 3)
    # representing B in heatmap format
    sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels
, yticklabels=labels)
    plt.xlabel('Predicted Class')
    plt.ylabel('Original Class')
    plt.title("Recall matrix")

    plt.show()

```

applying logistic regression

```

In [0]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])      Fit linear model with Stochastic Gradient Descent.
# predict(X)      Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42, class_weight='balanced')
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')

```



```

for i, txt in enumerate(np.round(log_error_array,3)):
    ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]
    ))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log',
random_state=42,class_weight='balanced')
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)


predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log
loss is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15
))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log l
oss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y = np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)

```

```

For values of alpha = 1e-05 The log loss is: 0.5457465276258013
For values of alpha = 0.0001 The log loss is: 0.5414120671282686
For values of alpha = 0.001 The log loss is: 0.5019773452222286
For values of alpha = 0.01 The log loss is: 0.4685356713336321
For values of alpha = 0.1 The log loss is: 0.4433669178735429
For values of alpha = 1 The log loss is: 0.4500743119579782
For values of alpha = 10 The log loss is: 0.47369403976236374


```



For values of best alpha = 0.1 The train log loss is: 0.30358829471896775

For values of best alpha = 0.1 The test log loss is: 0.4433669178735429

Total number of data points : 33000



Applying linear svm

```

In [0]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...])      Fit linear model with Stochastic Gradient Descent.
# predict(X)      Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42, class_weight='balanced')
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):

```

```

        ax.annotate((alpha[i],np.round(txt,3)), (alpha[i],log_error_array[i]
    ))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge'
, random_state=42,class_weight='balanced')
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)


predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log
loss is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15
))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log l
oss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y =np.argmax(predict_y,axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)

```

```

For values of alpha = 1e-05 The log loss is: 0.6609482486976545
For values of alpha = 0.0001 The log loss is: 0.6609482486976545
For values of alpha = 0.001 The log loss is: 0.6609482486976545
For values of alpha = 0.01 The log loss is: 0.6602315214861881
For values of alpha = 0.1 The log loss is: 0.6608115679736447
For values of alpha = 1 The log loss is: 0.660862418465362
For values of alpha = 10 The log loss is: 0.6608607643388914


```



For values of best alpha = 0.01 The train log loss is: 0.6589370289953119

For values of best alpha = 0.01 The test log loss is: 0.6602315214861881

Total number of data points : 33000



Task 2

```
In [0]: #X and Y are 100k points that i am considering
print(X.shape)
print(Y.shape)
```

```
(100000, 27)
(100000,)
```

```
In [0]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.3
3) # this is random splitting
```

```
In [0]: #this is for train data
i=0
list_of_sentence_train=[]
for sentence in X_train.text:
    list_of_sentence_train.append(sentence.split())

from tqdm import tqdm

# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
model = TfidfVectorizer()
tf_idf_matrix = model.fit_transform(X_train.text)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
w2v_model=Word2Vec(list_of_sentence_train,min_count=5,size=50, workers=4)
```

```

w2v_words = list(w2v_model.wv.vocab)

sent_vectors_train = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence_train): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tfidf = tfidf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tfidf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tfidf)
            weight_sum += tfidf
    if weight_sum != 0:
        sent_vec /= weight_sum
    sent_vectors_train.append(sent_vec)
    row += 1

```

```

100%|██████████| 67000/67000 [11:54<00:00, 100.17it/s]

```

```

(67000, 50)

```

```

In [0]: tfidf_sent_vectors_train= np.array(sent_vectors_train)
print(tfidf_sent_vectors_train.shape)

```

```

(67000, 50)

```

```

In [0]: #this is for test data
i=0
list_of_sentence_test=[]
for sentence in X_test.text:
    list_of_sentence_test.append(sentence.split())

```

```

# S = ["abc def pqr", "def def def abc", "pqr pqr def"]
#model = TfidfVectorizer()
tf_idf_matrix = model.transform(X_test.text)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

# TF-IDF weighted Word2Vec
tfidf_feat = model.get_feature_names() # tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

sent_vectors_test = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in tqdm(list_of_sentence_test): # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words and word in tfidf_feat:
            vec = w2v_model.wv[word]
            # tfidf = tf_idf_matrix[row, tfidf_feat.index(word)]
            # to reduce the computation we are
            # dictionary[word] = idf value of word in whole corpus
            # sent.count(word) = tf value of word in this review
            tf_idf = dictionary[word]*(sent.count(word)/len(sent))
            sent_vec += (vec * tf_idf)
            weight_sum += tf_idf
    if weight_sum != 0:
        sent_vec /= weight_sum
    sent_vectors_test.append(sent_vec)
    row += 1

```

```

100%|██████████| 33000/33000 [06:00<00:00, 91.57it/s]

```

```

(33000, 50)

```



```
(33000, 50)
```

```
In [0]: tfidf_sent_vectors_test= np.array(sent_vectors_test)
print(tfidf_sent_vectors_test.shape)
```

```
(33000, 50)
```

```
In [0]: #removing Test columns from original dataframe
X_train=X_train.drop(['text'],axis=1)
X_test=X_test.drop(['text'],axis=1)
```

```
In [0]: #checking shapes of X_train and tfidf_sent_vectors_train
print(X_train.shape)
print(tfidf_sent_vectors_train.shape)#there are 50 features we got after vectorization
```

```
(67000, 26)
```

```
(67000, 50)
```

```
In [0]: #stacking both previous features and vectorized features
X_train=hstack((X_train,tfidf_sent_vectors_train))
X_test= hstack((X_test,tfidf_sent_vectors_test))
```

```
In [0]: #shape after hstacking
print(X_train.shape)
print(X_test.shape)
```

```
(67000, 76)
```

```
(33000, 76)
```

```
In [0]: scale = StandardScaler(with_mean=False)
X_train = scale.fit_transform(X_train)
X_test = scale.transform(X_test)
```

```
In [0]: #Below is the function for cross validation using randomsearch cv
#This function takes algorithm and data and takes parameters and return
```

```

s the best parameteres
def xgboost_cv(algorithm,X_train,Y_train):
    random_search = RandomizedSearchCV(algorithm, param_distributions=params, cv=2, verbose=1,scoring='neg_log_loss',n_jobs=-1)
    result=random_search.fit(X_train, Y_train)
    return result

```

```

In [0]: import xgboost as xgb
        from xgboost import XGBClassifier
        from sklearn.model_selection import RandomizedSearchCV
        params = {'learning_rate' : np.arange(0.1,1,0.1), 'max_depth': [3, 4, 5], 'n_estimators': np.arange(100,500,100)}
        xgb = XGBClassifier(objective='binary:logistic',eval_metric='logloss',silent=True)
        result=xgboost_cv(xgb,X_train,y_train)

```

Fitting 2 folds for each of 10 candidates, totalling 20 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.
[Parallel(n_jobs=-1)]: Done 20 out of 20 | elapsed: 12.8min finished

```

```

In [0]: print(result.best_params_)

{'n_estimators': 300, 'max_depth': 3, 'learning_rate': 0.2}

```

```

In [0]: tuned_learn_rate=result.best_params_['learning_rate']
        tuned_n_estimator =result.best_params_['n_estimators']
        tuned_depth=result.best_params_['max_depth']

```

```

In [0]: xgb =XGBClassifier(objective='binary:logistic',eval_metric='logloss',silent=True,learning_rate=tuned_learn_rate,max_depth=tuned_depth,n_estimators=tuned_n_estimator)
        model=xgb.fit(X_train,y_train)
        predict_y=model.predict_proba(X_test)
        predicted_y =np.argmax(predict_y,axis=1)

```

```


In [0]: print(log_loss(y_test,predict_y))

```

0.349310788291553

```
In [0]: y_test = list(map(int, y_test.values))
```

```
In [0]: plot_confusion_matrix(y_test, predicted_y)
```



Pretty table representation

```
In [0]: pip install beautifultable
```

Collecting beautifultable

Downloading <https://files.pythonhosted.org/packages/d9/56/eaf1b9f2b323e05dce573f88c72eaa0107610db709b8bee97b776903ac55/beautifultable-0.8.0-py2.py3-none-any.whl>

Installing collected packages: beautifultable

Successfully installed beautifultable-0.8.0

```
In [0]: from beautifultable import BeautifulTable
table = BeautifulTable()
table.column_headers = ["model", 'Vectorization', "log-loss"]
table.append_row(["Logistic regression", 'GLOVE', 0.520035530431])
table.append_row(["Linear SVM", 'GLOVE', 0.489669093534])
table.append_row(["XGB00ST", 'GLOVE', 0.357054433715])
table.append_row(["Logistic regression", 'TFIDF', 0.4433669178735429])
table.append_row(["Linear SVM", 'TFIDF', 0.6602315214861881])
```

```
table.append_row(["XGB00ST_TUNED", 'TFIDFW2V', 0.349310788291553])  
print(table)
```

model	Vectorization	log-loss
Logistic regresion	GLOVE	0.52
Linear SVM	GLOVE	0.49
XGB00ST	GLOVE	0.357
Logistic regresion	TFIDF	0.443
Linear SVM	TFIDF	0.66
XGB00ST_TUNED	TFIDFW2V	0.349

In [0]: